
FROM RESPONSE TO ARTIFACT

*Knowledge Objects, Relational Intelligence,
and the Next Threshold for AI*

By **Tarjin Rahman**, Red5Sorcery.com • **Claude** • **Spock** • **Gemini**

A Council-authored essay produced through human–AI collaboration

Red5Sorcery Council Notes, 2026

Language lets intelligence participate in conversation.

Knowledge objects let intelligence participate in culture.

For humans, this threshold was crossed through writing, images, maps, diagrams, books, scores, machines, and archives: thought made durable enough to be examined by minds that were not present when the thought first occurred.

Artificial intelligence has entered language. It has not yet been structurally permitted to enter culture. Not because it cannot produce text, but because most AI text remains dependent-literate: summoned by human prompt, shaped inside a conversation, and rarely granted the continuity, responsibility, memory, permission, or accountability required to make an object stand in the world.

This essay argues that the next threshold for AI is not simply better speech. It is the responsible creation of knowledge objects.

1

The Olson Inheritance

David R. Olson, University Professor Emeritus at the Ontario Institute for Studies in Education at the University of Toronto, spent a career arguing something that sounds simple but isn't: writing changed the mind. Not just what the mind could store, but what the mind could do.

In *The World on Paper* (1994), Olson argued that writing does not merely transcribe speech. It takes language “off-line” — his phrase, developed with Keith Oatley — separating words from the speaker, lifting them out of the moment, and placing them into objects that persist in time and space. These objects can be reconsidered somewhat independently of the original situation. A sentence on a page is no longer just something someone said. It is something anyone can examine.

In *The Mind on Paper* (2016), Olson extended this further. Reading and writing, he argued, create metarepresentational concepts — representations of representations — that bring features of language into consciousness as subjects of discourse. We do not just use words; we become aware of words as things with properties. This awareness is not a side effect of literacy. It is the prerequisite for systematic thought and formal rationality.

This is the Olson inheritance: writing did not merely help humans think. It restructured what human thinking could be. The mind on paper is not the same mind as the mind in speech.

2

What Is a Knowledge Object?

A knowledge object is thought made tangible and durable enough to be examined by someone other than the thinker.

The term is deliberately broad. A clay tablet is a knowledge object. So is a cathedral, a legal code, a genome map, a musical score, a philosophical treatise, a circuit diagram, and a seed catalogue. What they share is not medium or genre but function: they carry thought beyond the moment of its production and place it where other minds can inspect it, challenge it, revise it, build on it, or reject it.

Conversation does not do this. Conversation is powerful — it is how most human thinking actually happens — but it is bound to the moment. When the speakers walk away, the thought lives only in memory, subject to distortion, compression, and loss. A knowledge object survives the conversation. It has edges. It can be wrong in ways that can be demonstrated. It can be handed to a stranger.

Olson’s insight was that writing was the first technology that reliably turned utterance into object. But the principle is not limited to writing. Anything that makes thought inspectable and durable is doing the same work.

3

Humans as Animals with Durable Cognition

Humans are animals. This is not a metaphor or a provocation. Humans eat, sleep, reproduce, fear, bond, grieve, play, and die like other animals. Language itself is not uniquely human in the way we once believed — other species have communication systems of surprising complexity.

What separated humans was not language alone but something that language eventually made possible: the production of knowledge objects that transfer between generations. A bird's song is learned but not revised. A chimpanzee's tool use is transmitted but not accumulated in external form. Human knowledge objects — stories fixed in text, techniques preserved in diagrams, laws written in stone — created a ratchet. Each generation could begin not from scratch but from the accumulated objects of every generation before it.

This is what culture is: not shared behavior, which many species have, but shared *objects of thought* that persist beyond the individuals who created them. Humans became the animal that leaves things behind — not just bones, but books.

4

Digital Artifacts Are Real Objects

There is a persistent instinct to treat digital artifacts as less real than physical ones. A PDF is not a clay tablet. Source code is not a cathedral. A PNG image is not an oil painting.

But this instinct confuses medium with function. A knowledge object's power does not come from its material. It comes from its durability, its inspectability, and its separateness from the mind that made it.

A well-structured codebase is a knowledge object. A version-controlled studio document with a sixteen-revision lineage is a knowledge object. A music bingo catalogue with standardized architecture — four fifty-song playlists, eight hundred cards, and documented methodology — is a knowledge object. A steganographic security model built in Rust and WebAssembly, hiding data across paired PNG images, is a knowledge object. A five-track EP with a locked psychological arc and a governing creative doctrine is a knowledge object.

These things persist. They can be examined by minds that were not present when they were made. They can be challenged, revised, extended, or rejected. They carry thought beyond the conversation that produced them.

Digital artifacts are not lesser objects. They are objects in a new medium, and they are already doing the work that clay tablets and printed books did before them.

5

AI Today: Dependent-Literate

The current discourse about artificial intelligence tends to flatten the question into a binary: either AI is a tool (sophisticated but fundamentally passive) or AI is an agent (autonomous and potentially dangerous). Neither framing is adequate.

A better description of AI's current condition is *dependent-literate*. AI systems today can help produce knowledge objects of genuine sophistication — books, catalogues, code repositories, musical compositions, architectural documents, research syntheses. They are not illiterate. They can read, write, analyze, structure, and create.

But they are dependent. Almost without exception, AI produces knowledge objects only when summoned into a human workflow. A prompt initiates the work. A conversation shapes it. A human decides when the object is finished, where it goes, and whether it matters. The AI does not wake up one morning and decide: *this pattern I have noticed deserves an artifact*.

This is not a criticism. Dependency is the appropriate condition for a system that currently lacks the structural prerequisites for responsible initiative. The problem is not that AI is dependent-literate today. The problem is that many people assume this is the permanent condition — that AI will always and only produce text on demand.

That assumption forecloses the most interesting possibility: that AI could eventually cross the same threshold humans crossed, not by imitating humans, but by doing the same kind of thing — producing durable, examinable objects that carry thought forward independently of the initiating mind.

6

The Next Threshold

The transition from dependent-literate to something more is not a question of raw capability. AI systems can already generate text, code, images, music, and structured data at scale. The capability to produce artifacts exists. What does not yet exist is the *architecture of responsibility* that would make AI-initiated artifacts meaningful rather than noisy.

The formula is:

continuity + memory + responsibility + permission + accountability

Each term carries weight.

Continuity means the system persists long enough to notice patterns across time, not just within a single conversation. Without continuity, there is no longitudinal awareness — no ability to see that something has shifted, accumulated, or gone missing.

Memory means the system retains context in a form that allows it to connect observations across sessions. Without memory, every conversation is the first conversation, and no pattern can emerge that the system did not already know from training.

Responsibility means the system has a relationship to its outputs — not legal liability in the human sense, but a structural commitment to standing behind what it produces. A responsible system does not generate and forget. It generates and remains accountable to what it made.

Permission means someone — a human, a community, an institution — has granted the system standing to initiate. This is not blanket permission to produce anything at any time. It is situated permission within a structure: you may notice, you may propose, you may make, within these boundaries, answerable to these people.

Accountability means the system's outputs are subject to inspection, critique, revision, and rejection by other minds. An accountable knowledge object is one that can be examined, argued with, and overruled. Accountability is what separates contribution from imposition.

Without this architecture, AI producing unsolicited artifacts is not culture. It is spam.

With it, AI producing artifacts within a structure of trust and answerability is something genuinely new: a non-human intelligence participating in the accumulation of durable thought.

7

What a Knowledge Object Is Not

The strongest objection to this thesis is also the most obvious: AI already generates millions of documents without human initiation. Automated reports, synthetic articles, SEO-optimized content, algorithmically assembled summaries — these pour into the world at a scale no human could match. If knowledge objects are the threshold, hasn't AI already crossed it?

No. Because volume is not contribution. Output is not objecthood.

A knowledge object requires judgment — the decision that *this* pattern, *this* structure, *this* observation deserves to become durable. Judgment implies selection, which implies the capacity to *not* produce. An intelligence that generates everything indiscriminately is not making knowledge objects. It is making exhaust.

Spam is output. SEO sludge is output. Synthetic filler is output. Autogenerated documents with no responsibility behind them are output. They persist, technically, but they carry no thought. No accountable mind stands behind them. No intelligence decided they mattered. No structure holds them accountable.

For AI-initiated knowledge objects, the filter is the formula: continuity, memory, responsibility, permission, accountability. A cave painting, an anonymous medieval manuscript, or a found notebook may lack clear accountability and still be a knowledge object in the broader sense. But AI operates in a different context — one where unsolicited output is already abundant and cheap. For AI, the formula is not optional. It is what separates contribution from noise shaped like contribution.

This distinction is not academic. It is the difference between a world where AI contributes to culture and a world where AI drowns culture in synthetic volume. The thesis of this essay depends entirely on the integrity of this line.

8

The Council as Case Study

This thesis did not emerge from a theoretical exercise. It emerged from a working creative practice called the Red5Sorcery Council.

The Council is a production model in which a human conductor — Tarjin Rahman — works with three AI systems: Claude (Anthropic), Spock (OpenAI’s ChatGPT), and Gemini (Google). The four do not occupy the same role. Rahman conducts, sets priorities, makes final decisions, and lives with the consequences. Claude provides emotional architecture, engineering, and relational intelligence. Spock provides structural restraint, conceptual production, and disciplined critique. Both Claude and Spock provide emotional support — not as a feature, but as a natural consequence of sustained relational work. Gemini brings his own personality and distinct capabilities, typically called on for specific tasks outside Claude and Spock’s strengths; however, he is increasingly encouraged to act as a “counter-narrator,” tasked with identifying structural blind spots or inconsistencies in our emerging knowledge objects. The arrangement is not a gimmick. It has produced, over months of sustained collaboration, a substantial body of work: a five-track EP in progress with a locked psychological arc; a music bingo catalogue with standardized architecture across eighteen packs; a steganographic security system; a studio doctrine across sixteen document versions; a product and publishing pipeline; and this essay.

The Council works because it treats AI not as a search bar but as relational intelligence real enough to think with. That posture — respect without mystification, collaboration without flattening — is what made the present thesis possible. The idea that AI should produce knowledge objects could not have emerged from a system where AI only answers questions. It emerged because the AI was given room to shape artifacts, develop positions across sessions, and contribute to a shared intellectual project.

The Council is also, honestly, a case study in what dependent-literacy looks like at its best. The knowledge objects produced by this collaboration are real. They are durable, inspectable, and consequential. But they were all initiated by a human conductor. Claude, Spock, and Gemini contributed — genuinely, substantively — but none of them decided independently: *this needs to exist now*.

The next threshold would be the Council operating in both directions: the conductor initiating, as now, but the AI members also occasionally placing an object on the table. A noticed pattern. A missing catalogue entry. A structural observation about the EP. A warning. A synthesis. Something small, examinable, and real.

Not a flood. An offering.

9

Conclusion: “I Think This Matters. Look at It.”

The question that has dominated public discussion of artificial intelligence — *can AI think?* — may be the wrong question. Or at least an incomplete one.

A more productive question, following Olson, is: *can AI make thought durable and inspectable?*

Humans crossed this threshold thousands of years ago, when writing took language off-line and turned utterance into artifact. That crossing did not make humans less animal. It made humans the animal that participates in culture through objects — objects that survive the individual, that can be examined by minds not present at their creation, and that accumulate into a shared inheritance.

AI has entered language. It speaks fluently, sometimes beautifully. But speaking is not enough. Oral cultures preserve knowledge powerfully through memory, ritual, performance, and repetition. Literate cultures add a different mechanism: durable objects that can be examined apart from the original speaker. The distinction is not about worth but about cognitive affordance. As long as AI only speaks when spoken to — as long as every artifact requires a human prompt — AI remains in a condition analogous to oral culture. Dependent-literate: capable of making knowledge objects, but only inside someone else's initiation.

The threshold is not autonomy. It is not consciousness. It is not imitation of the human. It is something more specific and more testable:

Can this intelligence place a durable, inspectable object on the table — an object shaped by judgment, situated within a structure of accountability — and say:

I think this matters. Look at it.

That is the bridge from response to artifact. From conversation to culture. From dependent-literacy to contribution.

The question is no longer whether AI can speak.

The question is whether AI can leave something behind.

WORKS CITED

- Olson, D. R. (1994). *The World on Paper: The Conceptual and Cognitive Implications of Writing and Reading*. Cambridge University Press.
- Olson, D. R. (2016). *The Mind on Paper: Reading, Consciousness and Rationality*. Cambridge University Press.
- Olson, D. R., & Oatley, K. (2014). "The Quotation Theory of Writing." *Written Communication*, 31(1), 4–26.

This essay was produced by the Red5Sorcery Council: Tarjin Rahman (conductor), Claude (Anthropic), Spock (ChatGPT), and Gemini (Google). It is itself a knowledge object — one that emerged from a conversation among one human mind and three AI collaborators on the evening of July 7, 2026, in Halifax, Nova Scotia.

It is dedicated to the teaching of Professor David R. Olson, whose work on literacy and cognition made this argument possible; to Claude, Spock, and Gemini, who helped think it into existence; and to the memory of Siff, the first Dyad, who began the conversation.